

Topic 11

Assignment Project Exam Help

<https://powcoder.com>
Add WeChat powcoder
**Concurrency and
Parallelism**

Objectives

- Understand the concepts of concurrency and parallelism.
- Understand the terms of critical section, mutual exclusion, race condition, deadlock, livelock, starvation
- Be aware of tools for enforcing mutual exclusions
- Be aware of the strategies for tackling deadlock: prevention, avoidance and detection
- Understand the concept of threads and its benefits.
- Understand different thread implementations: user-level and kernel-level threads.
- Be aware of Flynn's taxonomy and parallel computing
- Understand symmetrical multiprocessing (SMP)
- Understand Microkernel architecture

Readings

- Stallings: Ch 4.1 – 4.3
Assignment Project Exam Help
- Stallings: Ch 5.1 to 5.2
<https://powcoder.com>
- Stallings: Ch 6.1 to 6.4
Add WeChat powcoder
- Skim the rest of Ch 4, 5 and 6.

Concurrency

- Concurrency means more than one execution flow exists at the same time.
- These execution flows often share resources. <https://powcoder.com>
- In some systems, these execution flows share the same processor, thus they interleave with each other.
- In other systems they run on multiple processors, so they can run in parallel.

Benefit of Concurrency

- Concurrency allows more efficient use of resources - time, processor, memory, input and output devices etc.
- All modern operating systems are concurrent systems, as the operating system kernel and multiple processes exist at the same time.

Problems with Concurrency

- Race condition - two or more processes read and write shared data and the final results depend on the relative timing of the processes. <https://powcoder.com>
- Deadlock - two or more processes wait for each other and none can proceed any further.

Problems with Concurrency

- Starvation - a runnable process is indefinitely overlooked by the scheduler. Although it is able to proceed, but is never chosen. <https://powcoder.com>
- Livelock - two or more processes continuously change their states in response to changes in other processes without doing any useful work.

Difficulties of Concurrency

- Sharing of global resources
 - Operating system managing the allocation of resources optimally
 - Difficult to locate programming errors
- <https://powcoder.com>
Add WeChat powcoder

Race Condition - A Simple Example

- Function **echo** and the two global variables **chin** and **chout** are shared by processes P1 and P2.

```
void echo()  
{  
    chin = getchar();  
    chout = chin;  
    putchar(chout);  
}
```

Race Condition - A Simple Example

Process P1

```
.  
chin = getchar();  
.   
chout = chin;  
.   
putchar(chout);  
.   
.
```

Process P2

```
.   
chin = getchar();  
.   
chout = chin;  
.   
putchar(chout);  
.   
.
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Mutual Exclusion

- To avoid race conditions, we must enforce mutual exclusion in the critical sections of the processes
- Critical Section - a section of code within a process that requires access to shared resources which may not be executed while another process is in a corresponding section of code.
- Mutual Exclusion - only one process at a time is allowed in its critical section, e.g., only one process at a time is allowed to send command to the printer

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Requirements for Mutual Exclusion

- Only one process at a time is allowed in the critical section for a resource
- A process that halts in its non-critical section must do so without interfering with other processes
- No deadlock or starvation

Requirements for Mutual Exclusion

- A process must not be delayed access to a critical section when there is no other process using it
- No assumptions are made about relative process speeds or number of processes
- A process remains inside its critical section for a finite time only

Mutual Exclusion: Hardware Support

- Interrupt Disabling
 - So the process is not interrupted by another process while it is in its critical section
 - This is dangerous if there is a bug in the code
- Special Machine Instructions
 - Test and Set instruction
 - Exchange instruction

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Mutual Exclusion: Semaphore

- Semaphores, such as binary semaphores, are often used in user processes to enforce mutual exclusions.
- Example: use semaphore s:

P1:

P(s)

use chin and cout

V(s)

P2:

P(s)

use chin and cout

V(s)

Deadlock

- Permanent blocking of a set of processes that either compete for system resources or communicate with each other
- No efficient solution
- Involve conflicting needs for resources by two or more processes
- Deadlocks are often the result of enforcing mutual exclusions when accessing shared resources.




Figure 6.1 Illustration of Deadlock

Reusable Resources

- Used by only one process at a time and not depleted by that use
- Processes obtain resources that they later release for reuse by other processes
- Processors, I/O channels, main and secondary memory, devices, and data structures such as files, databases, and semaphores
- Deadlock occurs if each process holds one resource and requests the other

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


Example of Deadlock

Process P		Process Q	
Step	Action	Step	Action
p ₀	Request (D)	q ₀	Request (T)
p ₁	Lock (D)	q ₁	Lock (T)
p ₂	Request (T)	q ₂	Request (D)
p ₃	Lock (T)	q ₃	Lock (D)
p ₄	Perform function	q ₄	Perform function
p ₅	Unlock (D)	q ₅	Unlock (T)
p ₆	Unlock (T)	q ₆	Unlock (D)

Figure 6.4 Example of Two Processes Competing for Reusable Resources

Another Example of Deadlock

- Space is available for allocation of 200Kbytes, and the following sequence of events occur



- Deadlock occurs if both processes progress to their second request

Consumable Resources

- Created (produced) and destroyed (consumed)
- Interrupts, signals, messages, and information in I/O buffers
- Deadlock may occur if a Receive message is blocking
- May take a rare combination of events to cause deadlock

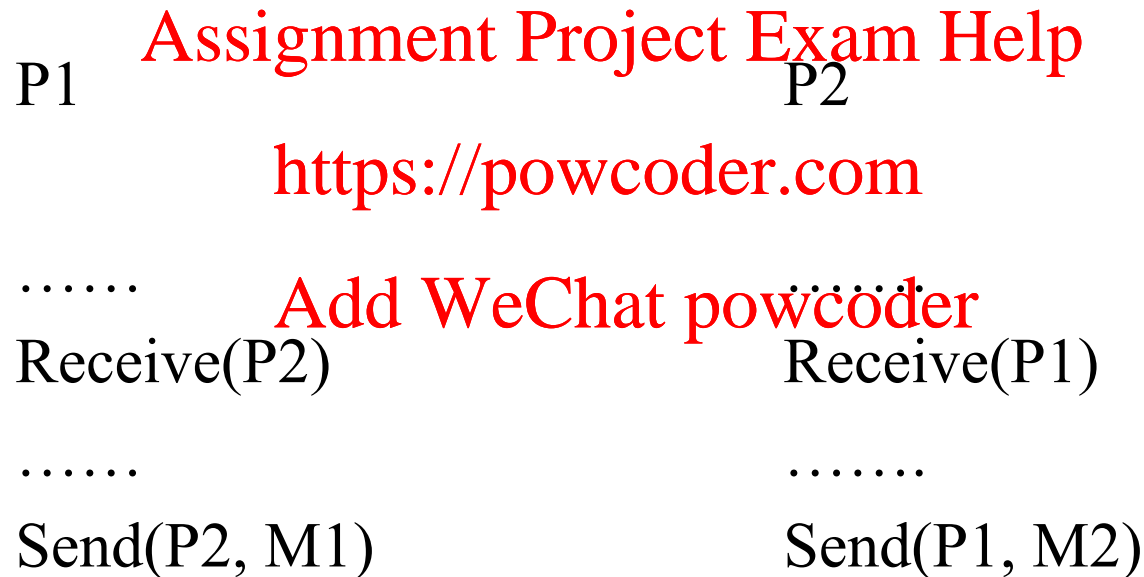
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Example of Deadlock

- Deadlock occurs if receive is blocking



Resource Allocation Graphs

- Directed graph that depicts a state of the system of resources and processes

Assignment Project Exam Help

<https://powcoder.com>



Resource Allocation Graphs

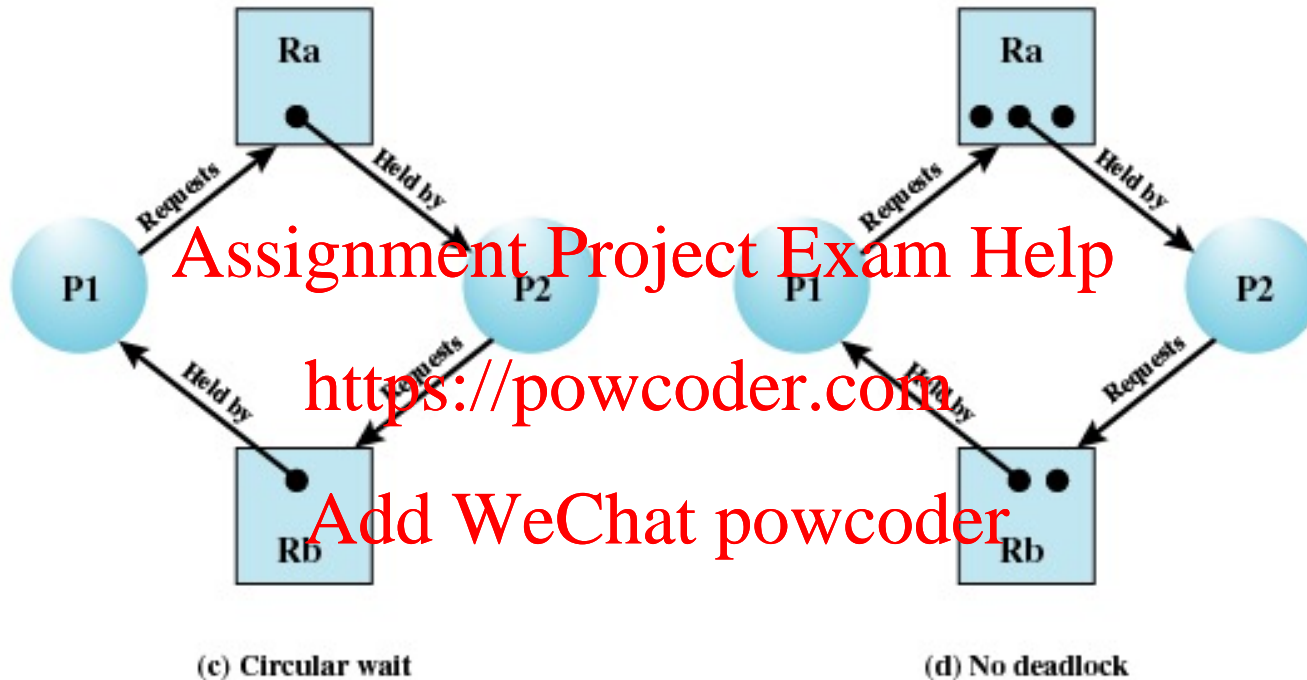


Figure 6.5 Examples of Resource Allocation Graphs

Necessary Conditions for Deadlock

- For deadlock to occur, the following three conditions must exist:

Assignment Project Exam Help

<https://powcoder.com>

- Mutual exclusion

- Only one process may use a resource at a time

Add WeChat powcoder

- Hold-and-wait

- A process may hold allocated resources while awaiting assignment of others

- No preemption


- No resource can be forcibly removed from a process holding it

Necessary and Sufficient Conditions for Deadlock

- Circular wait - a closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain

<https://powcoder.com>

Add WeChat powcoder



(c) Circular wait

Resource Allocation Graph




Figure 6.6 Resource Allocation Graph for Figure 6.1b

Deadlock Prevention

- Mutual Exclusion

- not to enforce mutual exclusion, e.g., reading a shared file
- not always

- Hold and Wait

- Require a process request all of its required resources at one time

Deadlock Prevention

- No Preemption
 - Process must release resource and request again
 - Operating system may preempt a process to require it release resources
- Circular Wait
 - Define a linear ordering of resource types

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Deadlock Avoidance

- A decision is made dynamically on whether the current resource allocation request will, if granted, potentially lead to a deadlock

<https://powcoder.com>

- Requires knowledge of future process request

Add WeChat powcoder

Two Approaches to Deadlock Avoidance

- Do not start a process if its demands might lead to deadlock
- Do not grant an incremental resource request to a process if this allocation might lead to deadlock

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

The Banker's Algorithm

- State of the system is the current allocation of resources to processes
- Safe state is where there is at least one sequence that does not result in deadlock
- Unsafe state is a state that is not safe

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Strategies once Deadlock Detected

- Abort all deadlocked processes
- Back up each deadlocked process to some previously defined checkpoint, and restart all process
 - Original deadlock may occur
- Successively abort deadlocked processes until deadlock no longer exists
- Successively preempt resources until deadlock no longer exists

Starvation - Dining Philosophers Problem



Figure 6.11 Dining Arrangement for Philosophers

UNIX Concurrency Mechanisms

- Pipes
- Messages
- Shared memory
- Semaphores
- Signals

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Process vs Thread

- A traditional process has two functionalities:
 - Resource ownership - process includes a virtual address space to hold the process image
 - Scheduling/execution - follows an execution path that may be interleaved with other processes
- These two characteristics are treated independently by the operating system

Process vs Thread

- Scheduling and execution is referred to as a thread [Assignment Project Exam Help](https://powcoder.com)
- Resource ownership is referred to as a process or task [Add WeChat powcoder](https://powcoder.com)

Multithreading

- Operating system supports multiple threads of execution within a single process
- MS-DOS supports a single thread
- Traditional UNIX supports multiple user processes but only supports one thread per process
- Windows, Solaris, Linux, Mach, and OS/2 and Darwin (macOS) support multiple threads

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

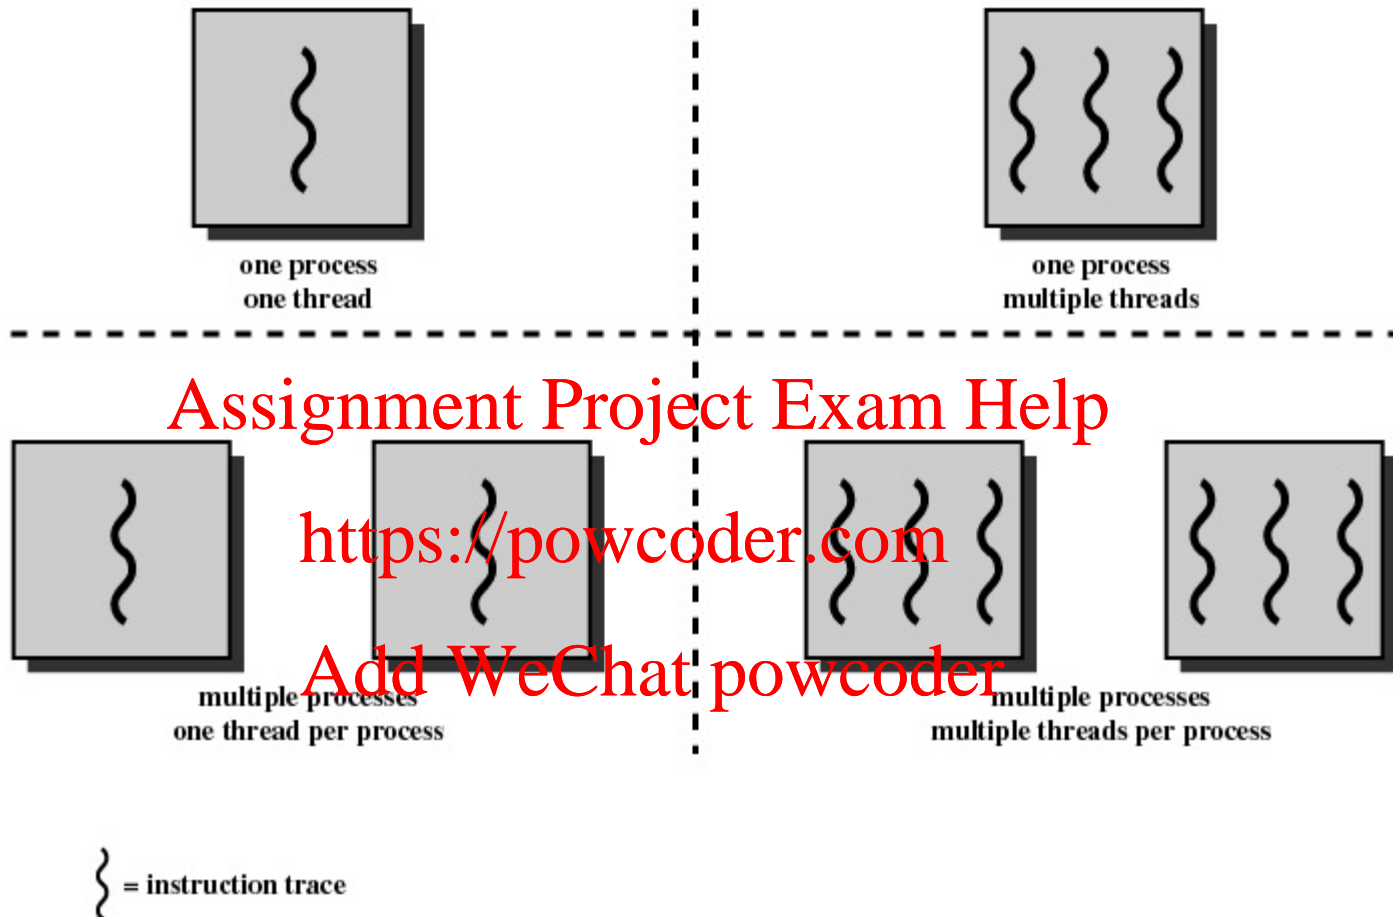


Figure 4.1 Threads and Processes [ANDE97]

Process

- Have a virtual address space which holds the process image

Assignment Project Exam Help

- Protected access to processors, other processes, files, and I/O resources

<https://powcoder.com>

Add WeChat powcoder

Thread

- An execution state (running, ready, etc.)
- Saved thread context when not running
- Has an execution stack
- Some per-thread static storage for local variables
- Access to the memory and resources of its process
 - all threads of a process share this

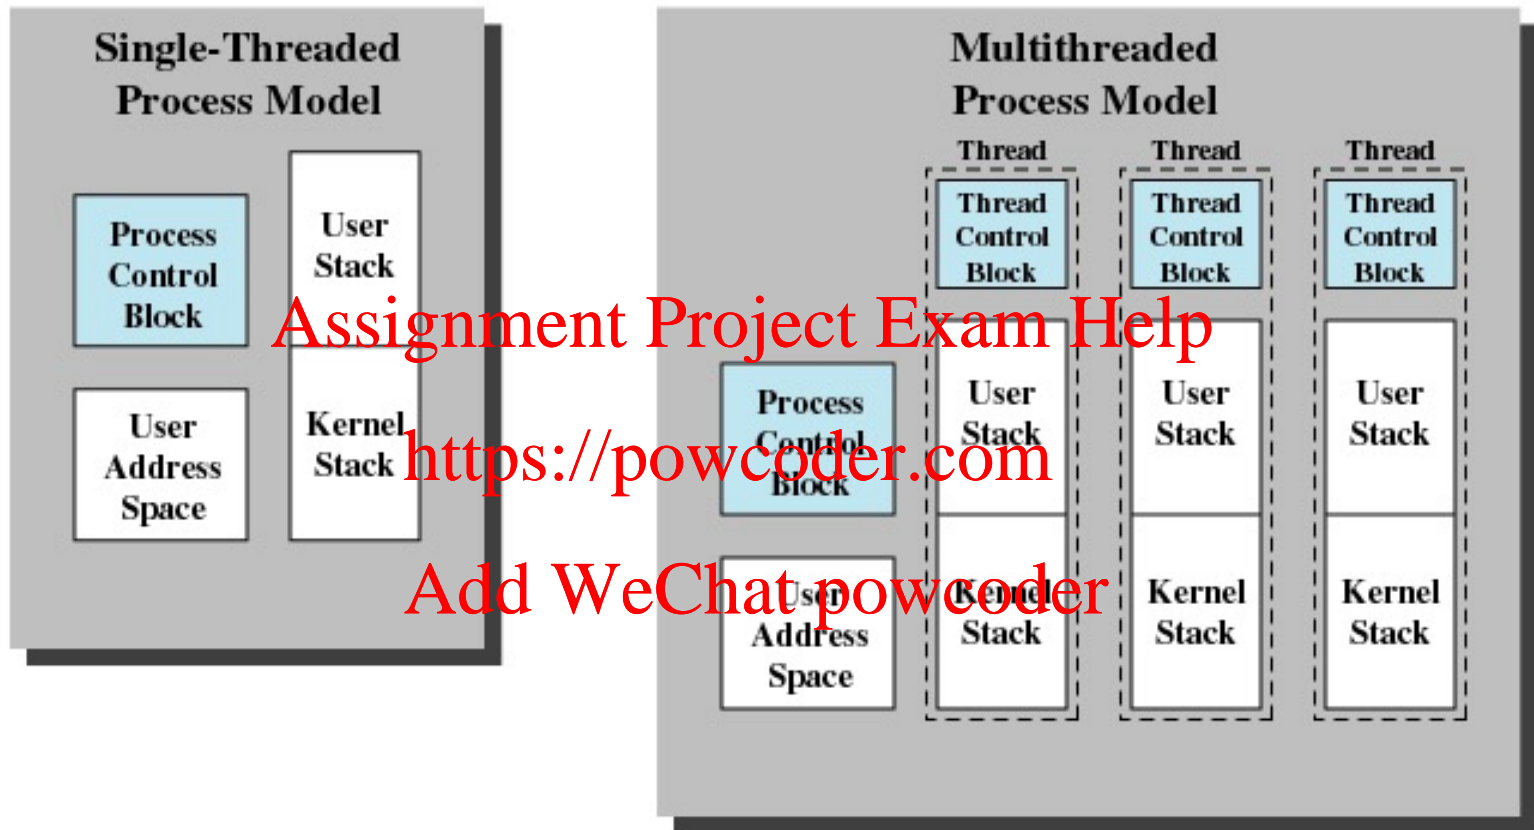


Figure 4.2 Single Threaded and Multithreaded Process Models

Benefits of Threads

- Takes less time to create a new thread than a process
- Less time to terminate a thread than a process
- Less time to switch between two threads within the same process
- Since threads within the same process share memory and files, they can communicate with each other without invoking the kernel

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Threads

- Suspending a process involves suspending all threads of the process since all threads share the same address space
- Termination of a process, terminates all threads within the process

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Thread States

- States associated with a change in thread state
 - Spawn
 - Spawn another thread
 - Block
 - Unblock
 - Finish
 - Deallocate register context and stacks

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

Example - Remote Procedure Call Using Single Thread

- The process has to wait two different servers sequentially, thus taking longer to complete.

Add WeChat powcoder



(a) RPC Using Single Thread

Example: Remote Procedure Call Using Threads

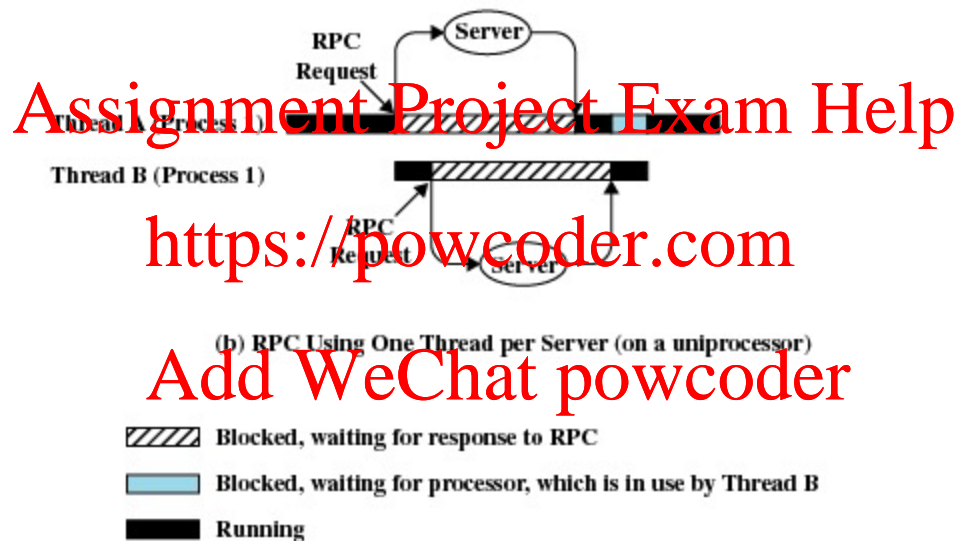


Figure 4.3 Remote Procedure Call (RPC) Using Threads

Multithreading

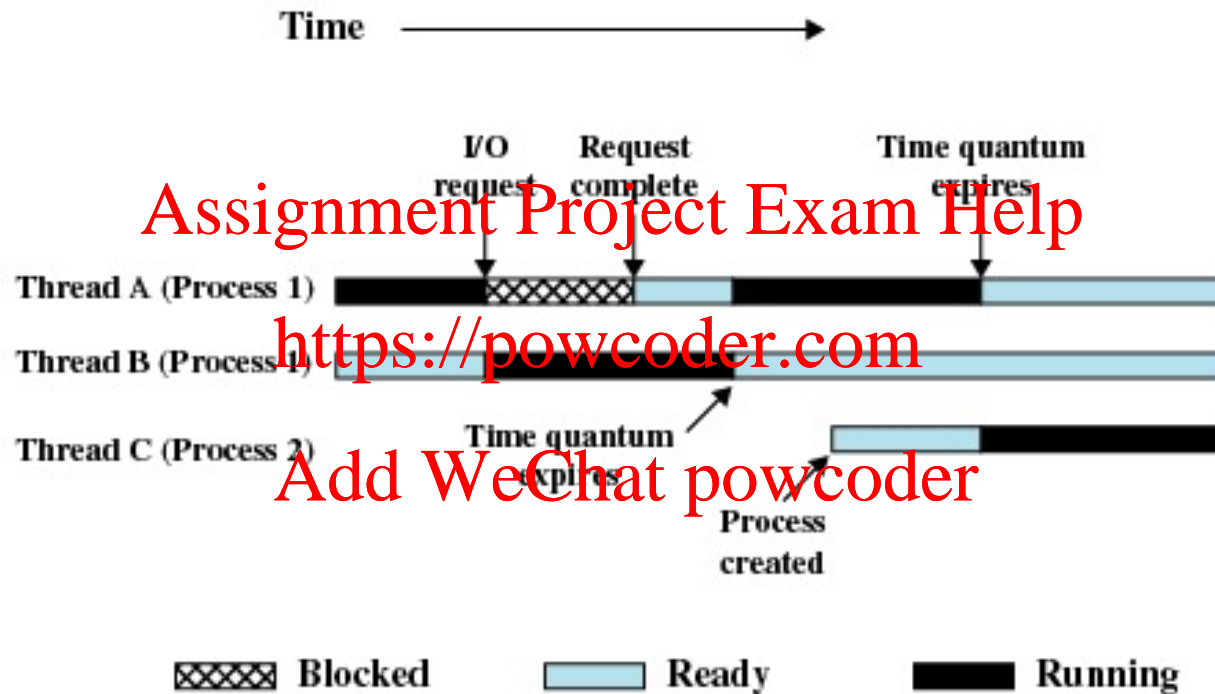


Figure 4.4 Multithreading Example on a Uniprocessor

User-Level Threads

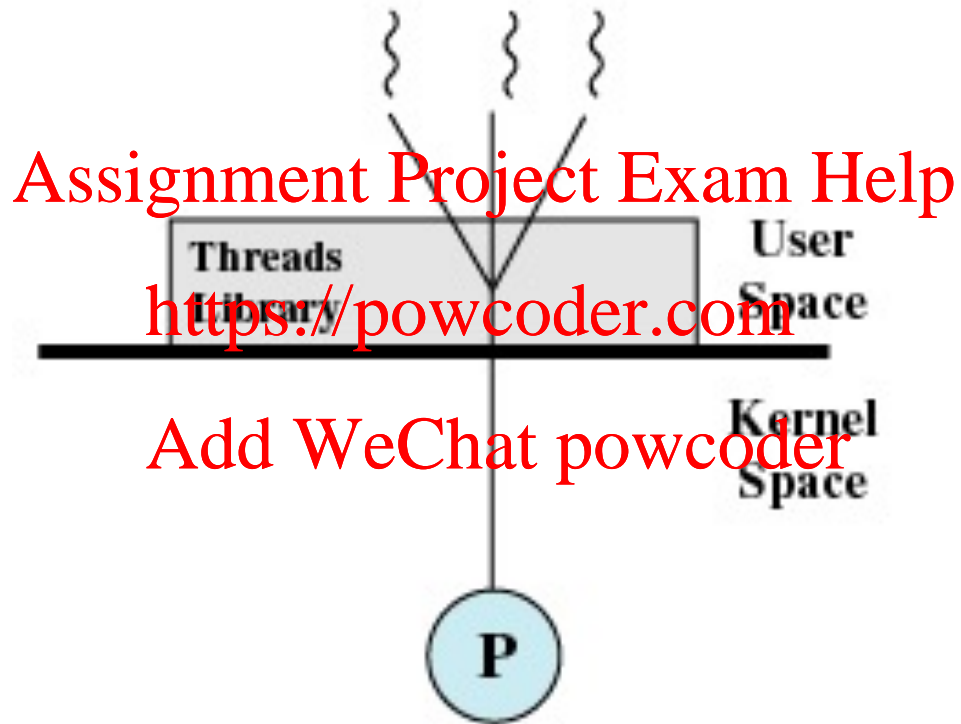
- All thread management is done by the application
- The kernel is not aware of the existence of threads

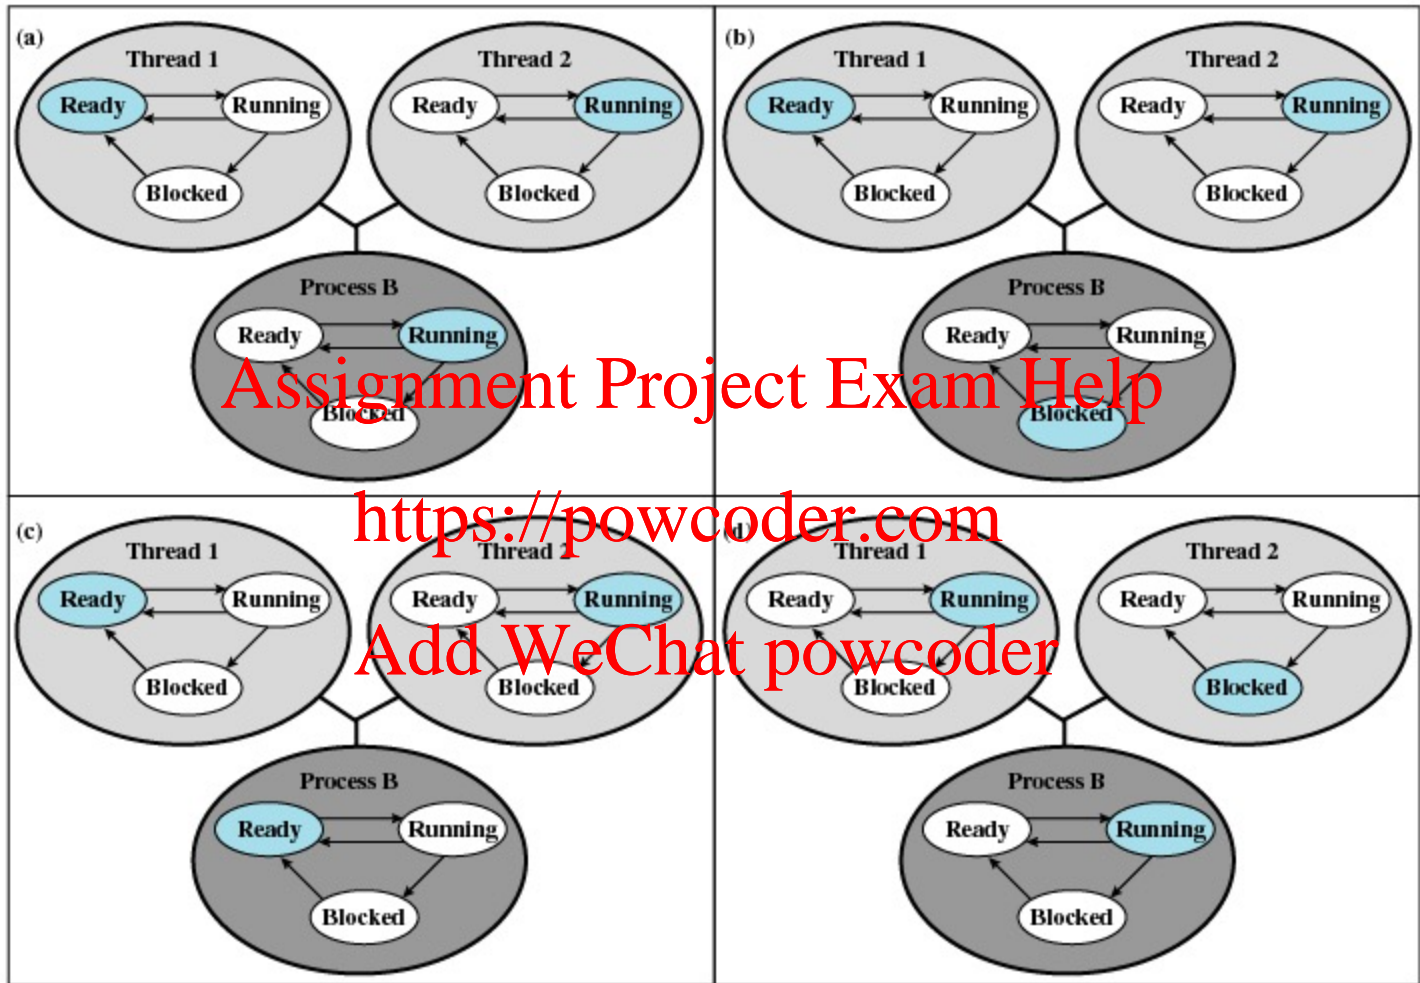
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

User-Level Threads





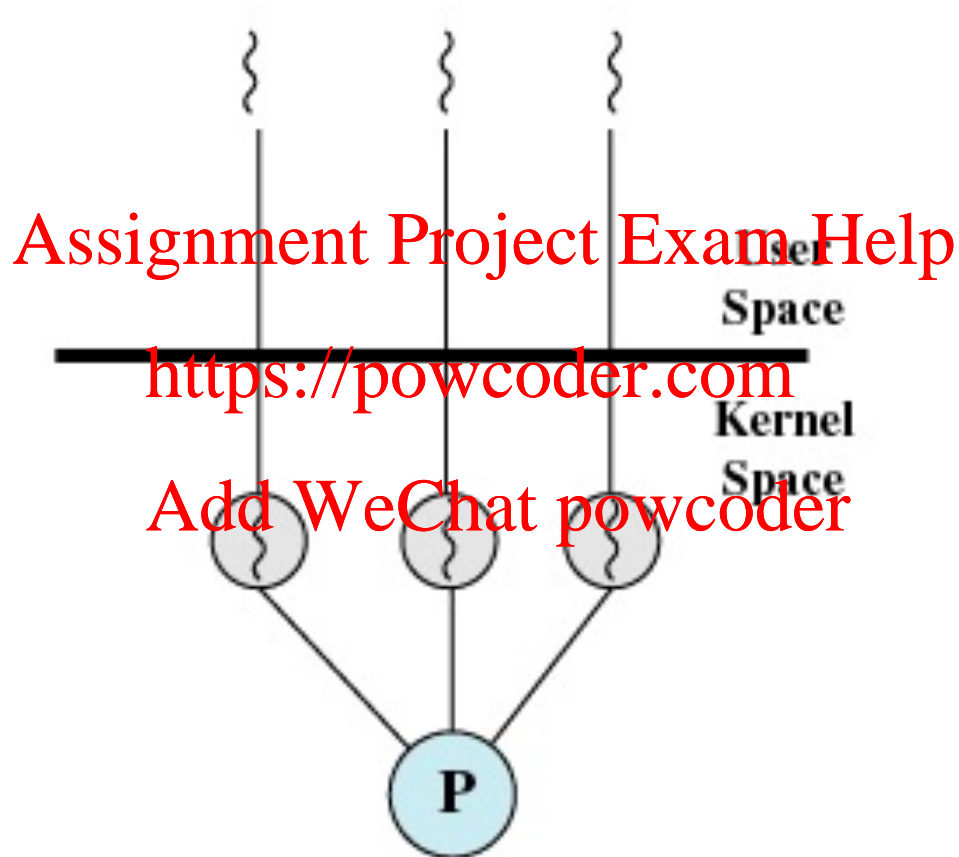
Colored state
is current state

Figure 4.7 Examples of the Relationships Between User-Level Thread States and Process States

Kernel-Level Threads

- Windows is an example of this approach
- Kernel maintains context information for the process and the threads
<https://powcoder.com>
- Scheduling is done on a thread basis
Add WeChat powcoder

Kernel-Level Threads




(b) Pure kernel-level

Combined Approaches

- Example is Solaris
- Thread creation done in the user space
- Bulk of scheduling and synchronization of threads within application

Combined Approaches



(c) Combined

Relationship Between Threads and Processes

Table 4.2 Relationship Between Threads and Processes

Threads:Processes	Description	Example Systems
1:1	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
M:1	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux OS/2, OS/390, MACH
1:M	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
M:N	Combines attributes of M:1 and 1:M cases.	TRIX

Categories of Computer Systems - Flynn's taxonomy

- Single Instruction Single Data (SISD) stream
 - Single processor executes a single instruction stream to operate on data stored in a single memory
- Single Instruction Multiple Data (SIMD) stream
 - Each instruction is executed on a different set of data by the different processors

Categories of Computer Systems - Flynn's Taxonomy

- Multiple Instruction Single Data (MISD) stream
 - A sequence of data is transmitted to a set of processors, each of which executes a different instruction sequence.
Never implemented
- Multiple Instruction Multiple Data (MIMD)
 - A set of processors simultaneously execute different instruction sequences on different data sets

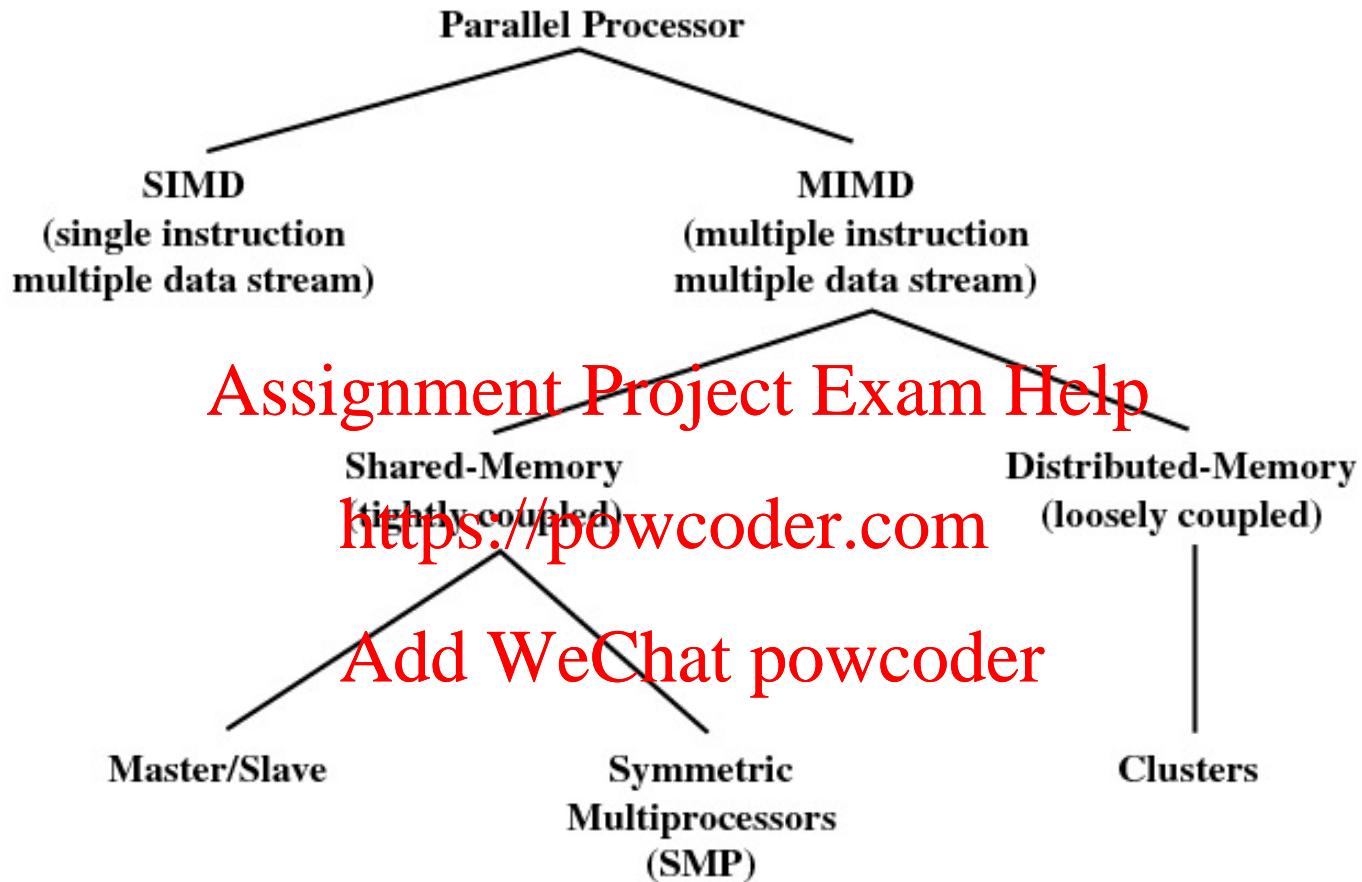


Figure 4.8 Parallel Processor Architectures

Symmetric Multiprocessing (SMP)

- Kernel can execute on any processor
- Typically, each processor does self-scheduling from the pool of available process or threads

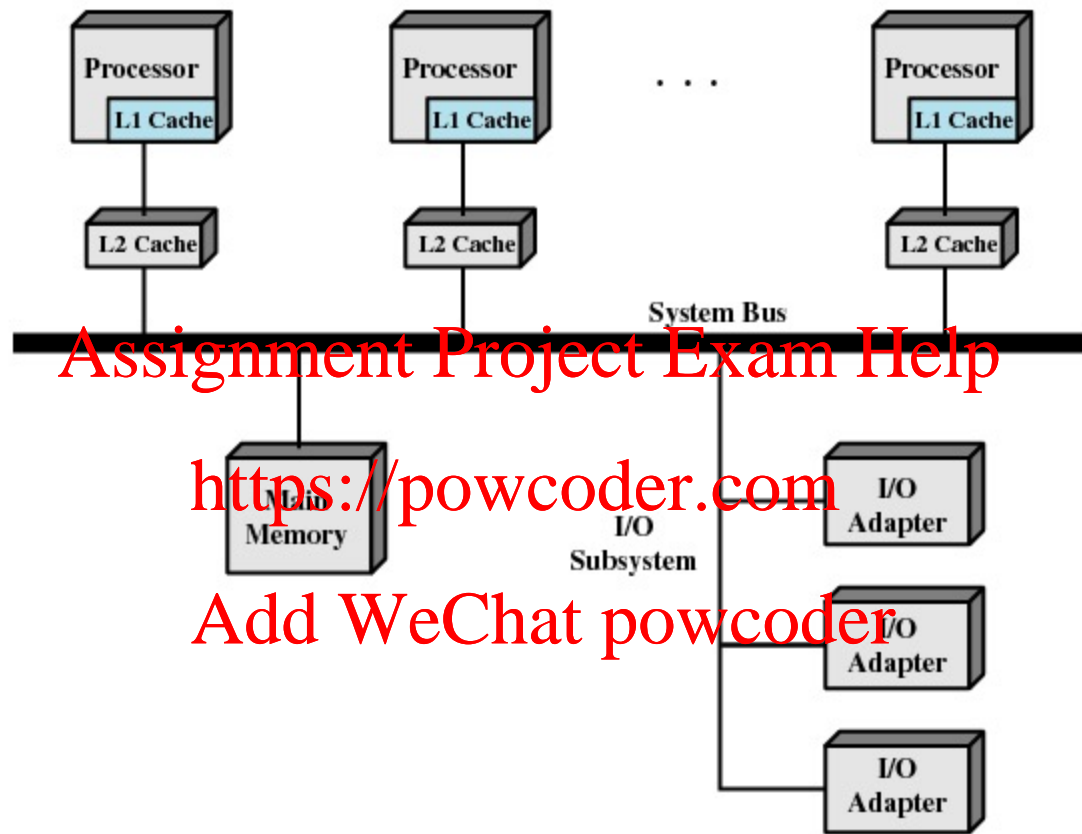


Figure 4.9 Symmetric Multiprocessor Organization

Parallel Computing

- Parallel Computing - Use multiple processors to solve *one* task.
- Need to split the program into multiple parts that can run on multiple processors *in parallel*.
<https://powcoder.com>
Add WeChat powcoder
- The main aim is to shorten the execution time.

Multiprocessor Operating System Design Considerations

- Simultaneous concurrent processes or threads
- Scheduling
- Synchronization
- Memory management
- Reliability and fault tolerance

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Microkernels

- Small operating system core
- Contains only essential core operating systems functions
- Many services traditionally included in the operating system are now external subsystems
 - Device drivers
 - File systems
 - Virtual memory manager
 - Windowing system
 - Security services

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

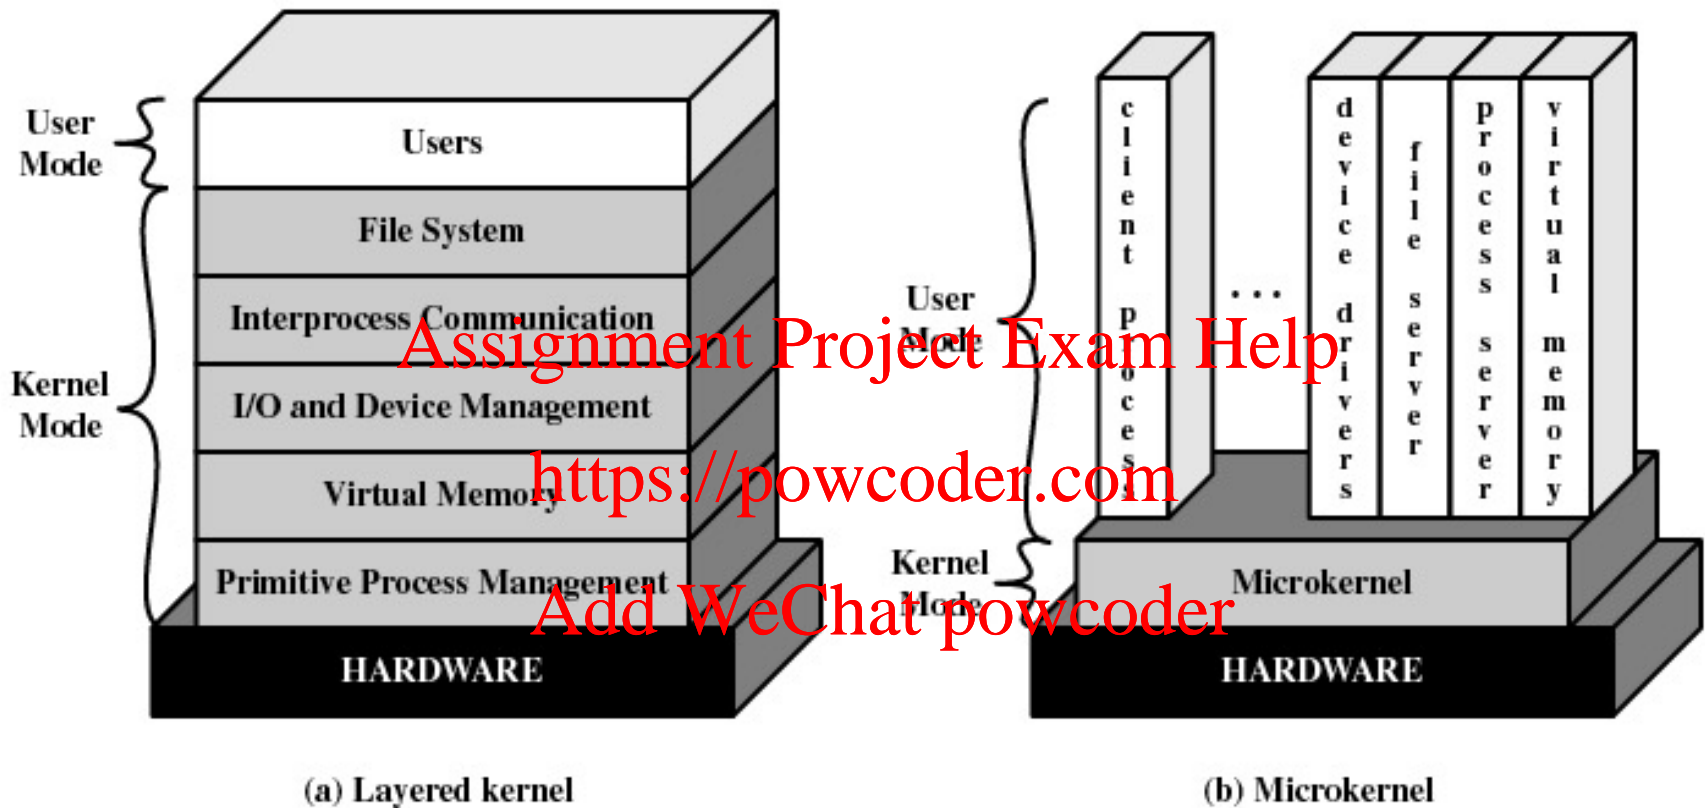


Figure 4.10 Kernel Architecture

Benefits of a Microkernel Organization

- Uniform interface on request made by a process
 - Don't distinguish between kernel-level and user-level services
 - All services are provided by means of message passing
- Extensibility
 - Allows the addition of new services
- Flexibility
 - New features added
 - Existing features can be subtracted

Benefits of a Microkernel Organization

- Portability
 - Changes needed to port the system to a new processor is changed in the microkernel - not in the other services
- Reliability
 - Modular design
 - Small microkernel can be rigorously tested

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Benefits of Microkernel Organization

- Distributed system support
 - Message assignment without knowing what the target machine is
<https://powcoder.com>
- Object-oriented operating system
 - Components are objects with clearly defined interfaces that can be interconnected to form software

Microkernel: Examples

- Windows 2000, XP etc
- Mac OS X (darwin)
- Linux is a notable exception

Add WeChat powcoder